

# Dynamic Data Race Detection

Aditya Thakur

# Data race

Two threads (goroutines) concurrently access a shared memory location and at least one of the accesses is a Write (i.e. updates the value).

# Data race - Example

```
var g int

func update() {
    g = g + 2
}

func main() {
    go update()
    go update()
}
```

# Using mutex to synchronize

```
var g int
var mu sync.Mutex

func update() {
    mu.Lock()
    g = g + 2
    mu.Unlock()
}

func main() {
    go update()
    go update()
}
```

# Using channels to synchronize

```
var g int
var ch chan int

func update() {
    ch <- 1
    g = g + 2
    <- ch
}

func main() {
    ch = make(chan int, 1)
    go update()
    go update()
}
```

# Impact of data races

- Can be harmless
  - Do not lead to any real errors in the program
- Can cause the program to crash, produce incorrect answers, or worse
  - Data race in Therac-25 medical electron accelerator lead to loss of lives
  - Northeast blackout of 2003 caused by data race

# Data race detection

Can we automatically detect data races in programs?

# Two approaches

- Static program analysis
  - Automated techniques to analyze the source code
  - Akin to finding bugs by staring at the source
  - Able to prove the absence of bugs; reported bugs might be false positives
- Dynamic program analysis
  - Automated techniques to analyze trace of events generated by executing the code
  - Akin to finding bugs by inserting print statements
  - Unable to prove the absence of bugs; reports fewer false positives



# Dynamic data race detection

Identify races by observing the **trace of events**

For each goroutine, sequence of

- reads/writes for each memory location
- locks/unlocks for each mutex
- Sends/receives for each channel

# Dynamic data race detection

- Two approaches
  - Happens-before analysis
  - Lockset analysis
- The go lang race detector uses a hybrid approach combining happens-before and lockset analysis


# Happens-before analysis

- Happens-before is a partial order of all events of all threads in a concurrent execution.
- **Event X happens-before event Y** if
  - X has been observed before Y,  
and
    - Events X and Y are in the same thread, OR
    - X is an Unlock() and Y is a Lock() for the same mutex OR
    - X is a send and Y is a receive for the same channel OR
    - There exists an event U such that X happens-before U, U happens-before Y

# Data races and Happens-before

If two threads (or goroutines) access a shared variables, and the accesses are *not* ordered by the *happens-before relation*, then a data race **could have occurred**.

# Detecting data race using happens-before

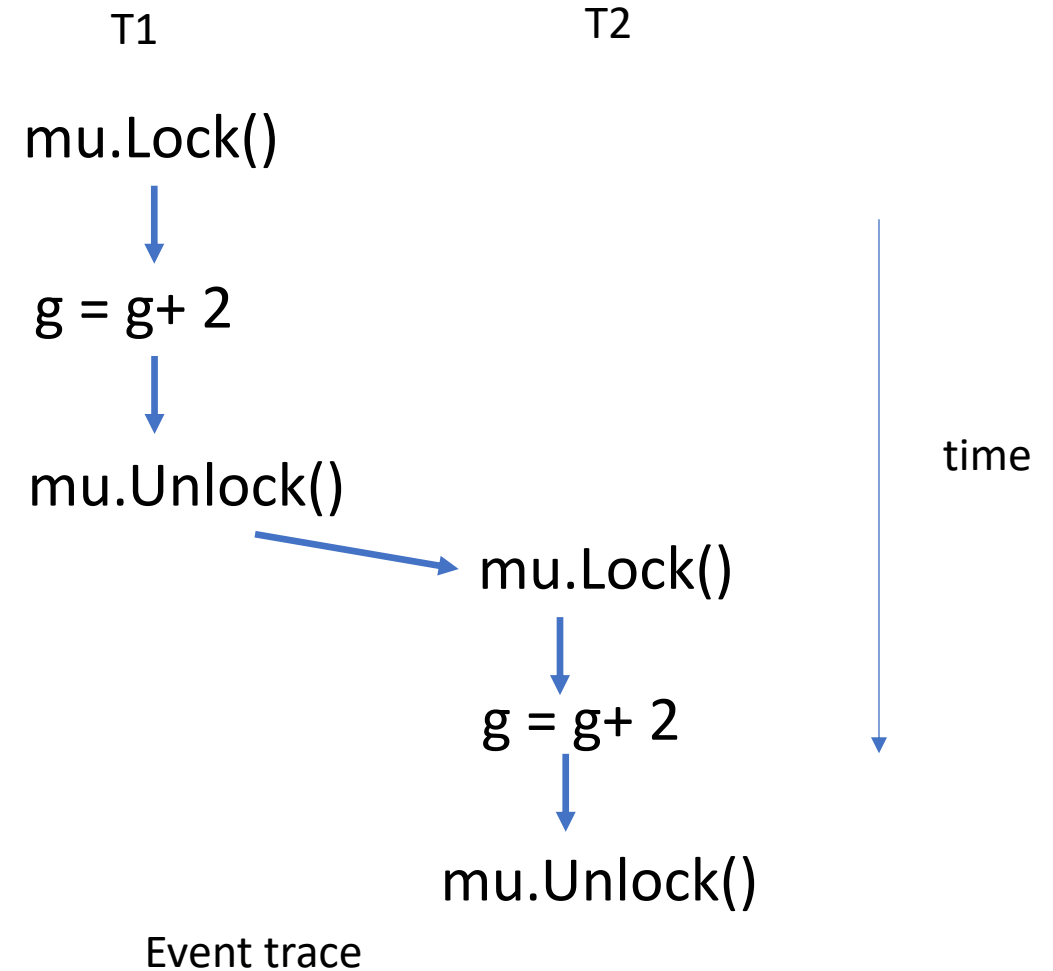


# Using mutex to synchronize

```
var g int
var mu sync.Mutex

func update() {
    mu.Lock()
    g = g + 2
    mu.Unlock()
}

func main() {
    go update()
    go update()
}
```

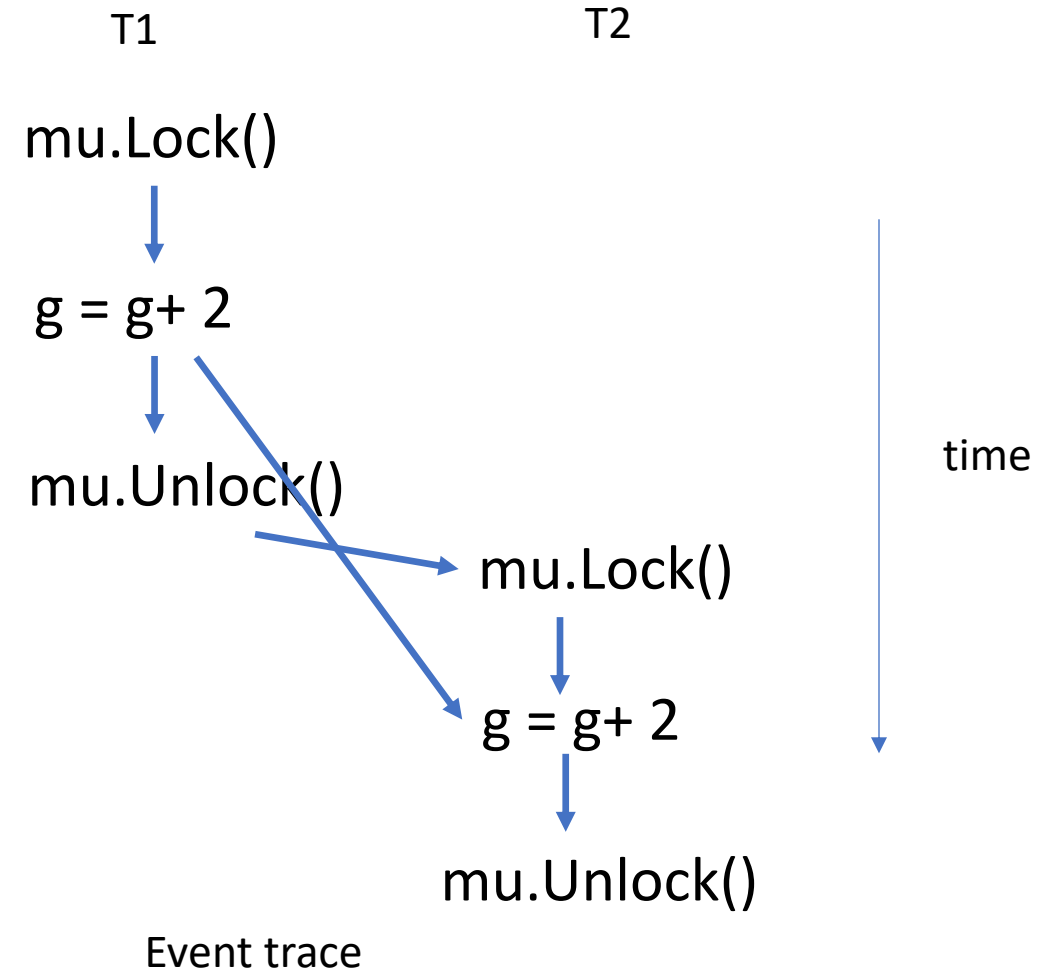


# Using mutex to synchronize

```
var g int
var mu sync.Mutex

func update() {
    mu.Lock()
    g = g + 2
    mu.Unlock()
}

func main() {
    go update()
    go update()
}
```




# Using channels to synchronize

```
var g int
var ch chan int

func update() {
    ch <- 1
    g = g + 2
    <- ch
}

func main() {
    ch = make(chan int, 1)
    go update()
    go update()
}
```






# Using channels to synchronize

```
var g int
var ch chan int

func update() {
    ch <- 1
    g = g + 2
    <- ch
}

func main() {
    ch = make(chan int, 1)
    go update()
    go update()
}
```



```

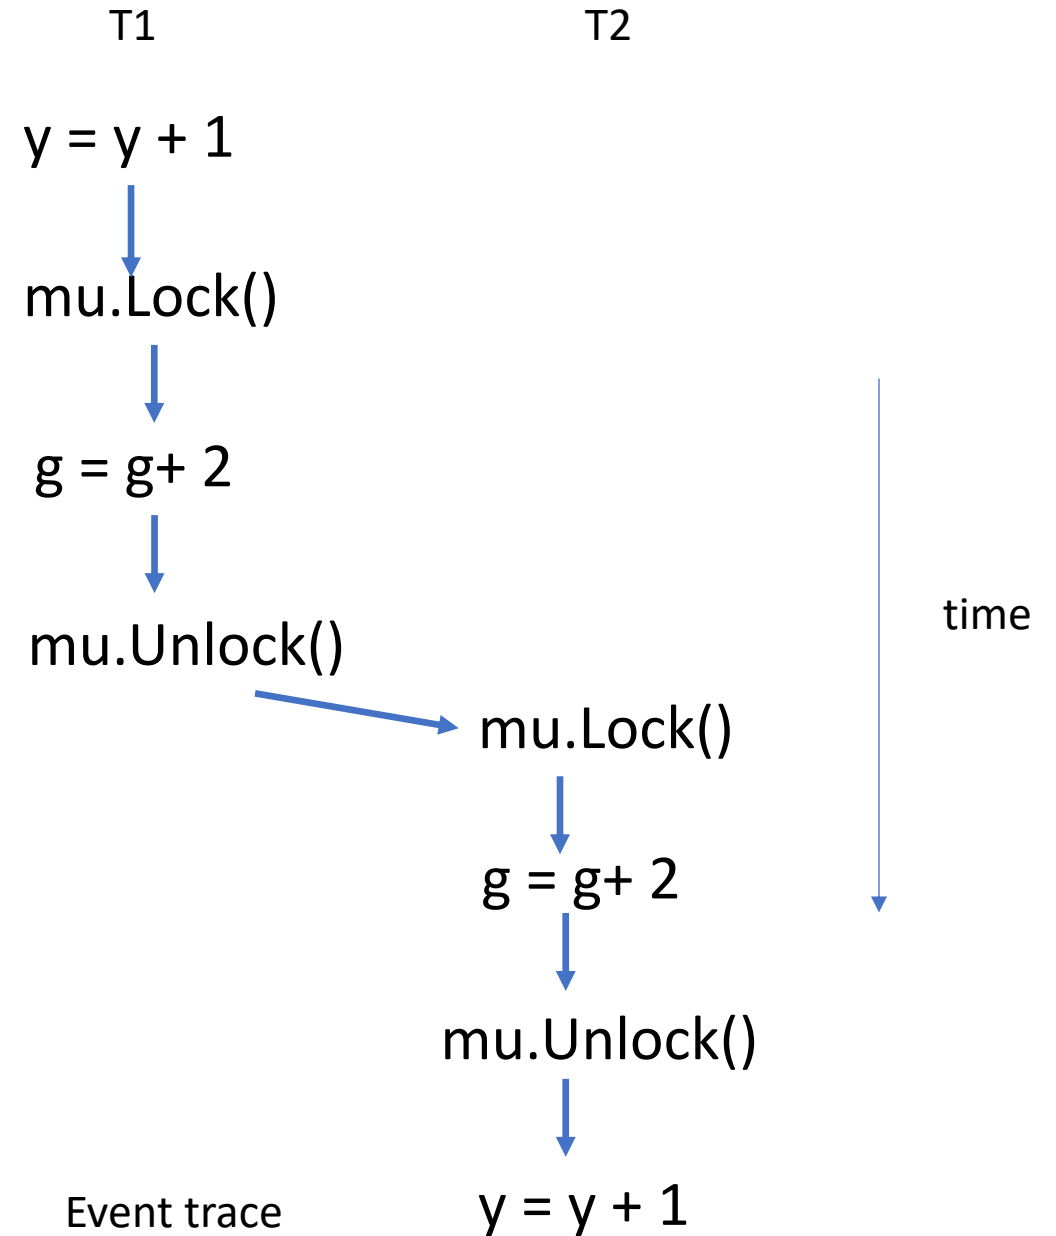
func update1() {
    y = y + 1
    mu.Lock()
    g = g + 2
    mu.Unlock()
}

var g, y int
var mu sync.Mutex

func update2() {
    mu.Lock()
    g = g + 2
    mu.Unlock()
    y = y + 1
}

func main() {
    go update1()
    go update2()
}

```




```
func update1() {
    y = y + 1
    mu.Lock()
    g = g + 2
    mu.Unlock()
}
```

```
var g, y int
var mu sync.Mutex
```

```
func update2() {
    mu.Lock()
    g = g + 2
    mu.Unlock()
    y = y + 1
}
```

```
func main() {
    go update1()
    go update2()
}
```



Event trace

# Lockset Analysis

- Keep track of which locks are used when accessing shared variables
- Report warning if no lock is consistently used when accessing a shared variable  $v$

# Lockset Algorithm

- Let  $locks\_held(t)$  be the set of locks held by thread  $t$
- For each  $v$ , initialize  $C(v)$  to the set of all locks.
- On each access to  $v$  by thread  $t$ ,
  - $C(v) := C(v) \cap locks\_held(t)$
  - If  $C(v) = \emptyset$ , then issue warning

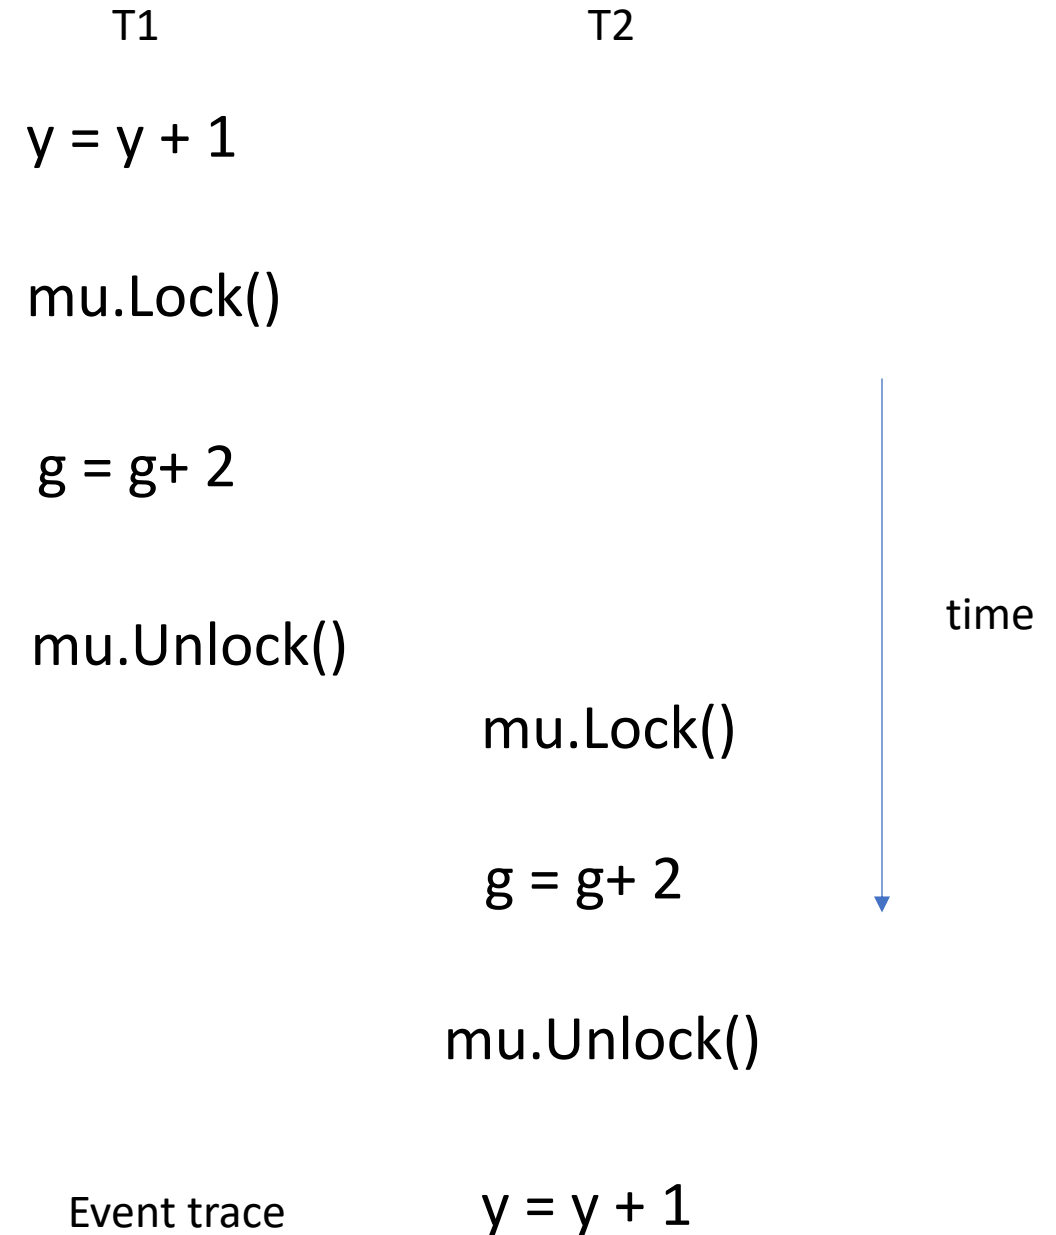
# What are the locksets for g and y ?

```
var g, y int
var mu sync.Mutex

func update1() {
    y = y + 1
    mu.Lock()
    g = g + 2
    mu.Unlock()
}

func update2() {
    mu.Lock()
    g = g + 2
    mu.Unlock()
    y = y + 1
}

func main() {
    go update1()
    go update2()
}
```



# What is the lockset for g?

```
var g int
var mu1, mu2, mu3 sync.Mutex
```

```
func update1() {
    mu1.Lock()
    mu2.Lock()
    g = g + 2
    mu2.Unlock()
    mu1.Unlock()
}
```

```
func update2() {
    mu2.Lock()
    mu3.Lock()
    g = g + 2
    mu3.Unlock()
    mu2.Unlock()
}
```

```
func update3() {
    mu1.Lock()
    mu3.Lock()
    g = g + 2
    mu3.Unlock()
    mu1.Unlock()
}
```

```
func main() {

    go update1()
    go update2()
    go update3()

}
```

Is there a race?

# What about channels?

```
var g int
var ch chan int

func update() {
    ch <- 1
    g = g + 2
    <- ch
}

func main() {
    ch = make(chan int, 1)
    go update()
    go update()
}
```

T1

ch <- 1

g = g + 2

<- ch

ch <- 1

g = g + 2

<- ch

time





# Golang data race detector

- The golang race detector uses a hybrid approach combining happens-before and lockset analysis

# References

- Lamport, Leslie. "Time, clocks, and the ordering of events in a distributed system." *Communications of the ACM* 21.7 (1978): 558-565.
- Savage, Stefan, et al. "Eraser: A dynamic data race detector for multithreaded programs." *ACM Transactions on Computer Systems (TOCS)* 15.4 (1997): 391-411.
- Serebryany, Konstantin, and Timur Iskhodzhanov. "ThreadSanitizer: data race detection in practice." *Proceedings of the Workshop on Binary Instrumentation and Applications*. ACM, 2009.