

Announcements

- Office hours today are 2:10pm–3:00pm in 2203 Watershed Sciences
 - That's the regular office hours shifted half an hour earlier

Memory Management

Paging

- Program memory need not be contiguous
 - Solves problem of compaction in MVT
- How it works
 - Virtual address split in two
 - High bits represent page number
 - Low bits represent page offset
 - Page table has the address of each page frame in physical memory

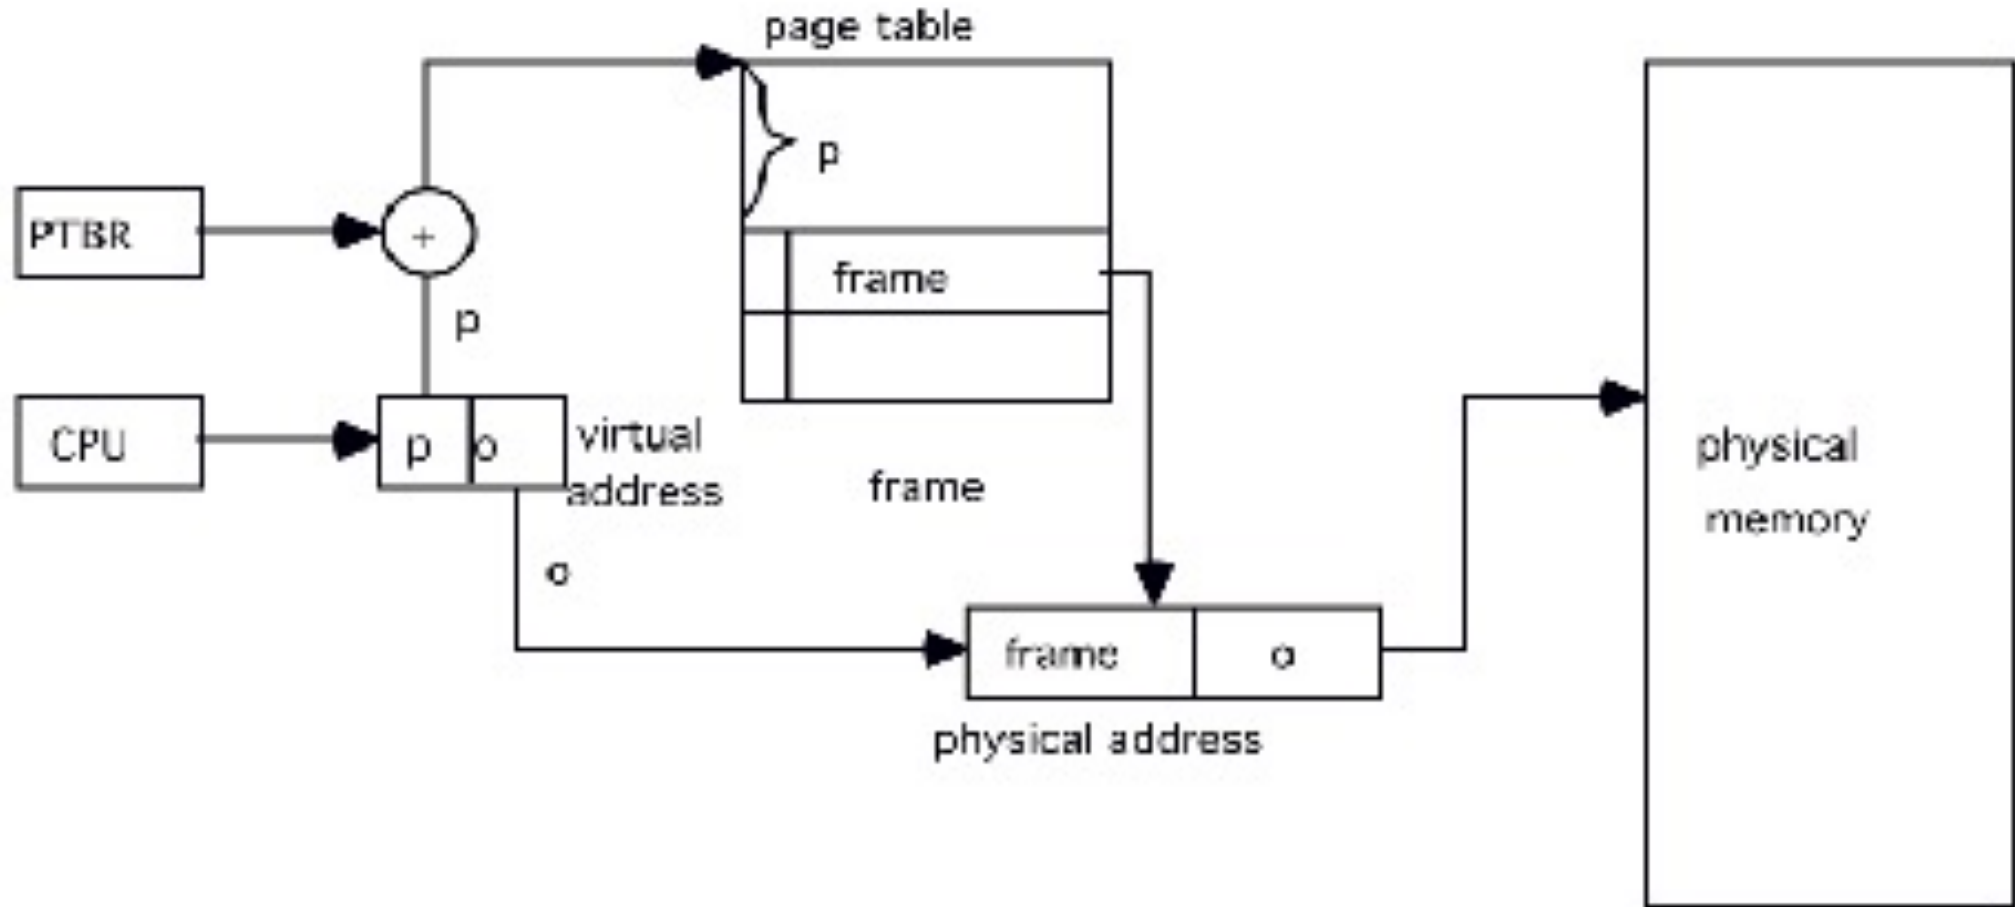
Page Frames

- Frame is physical memory into which a page is put
- Page is unit of virtual memory put into physical memory
- Both are same size, defined by hardware
 - Usually 1024, 2048, 4096, or 8192 words or bytes per page
- If page contains p words, virtual address va produces:
 - page number = va / p
 - page offset = $va \% p$
- If p is power of 2, then
 - page number = high order bits of va
 - page offset = low order bits of va

Address Translation

- Address of page table stored in Page Table Base Register (PTBR)
- Add this to page table number to get address of page table entry
- Get physical address of frame
- Add offset to that address to get physical memory address corresponding to virtual address

Address Translation



Example

- Virtual address is 7
- Page size is 2 words (not realistic)
- So page number: $7 / 2 = 3$
- And page offset: $7 \% 2 = 1$
- Frame is frame number 5
- Frame size = page size = 2, so base of frame is $5 \times 2 = 10$
- Offset is 1, so physical address of virtual word at 7 is $10 + 1 = 11$

page number	frame number
0	4
1	3
2	1
3	5
4	2
5	6

Process Scheduling

- Process memory size given in pages
- Process has n pages, so it needs n frames
- External fragmentation: *none*
- Internal fragmentation: at most $p-1$, where p is page size
 - So expected internal fragmentation per process is $p/2$

Storing Page Tables

- Small number of pages: use register for the page table
 - Loading, modifying these are privileged
 - Address translation very efficient, as registers use high speed logic
- Large number of pages: store page table in memory
 - For this (usual) case, you need the PTBR

Context Switching

- Page tables means changing 1 register, the PTBR
 - Value of swapped out process goes into the process' PCB
 - When swapped in, the PTBR is loaded with saved value

Memory Access Time

- Problem: now memory references are twice as slow!
 - First memory access is to the page table, to get physical address associated with virtual page
 - Next memory access is to the desired address
- Effective memory access time (EMAT)
 - Actual time needed to access memory

Optimizations

- Use *cache (associative memory, look-aside memory)*
- Registers store (key, value) pair
- Given key, cache hardware compares key to stored keys at once, returning corresponding values
- But this memory is expensive!

Caches and Paging

- Put some page table entries into cache
 - Usually too many to put them all there
- Here's what happens:
 - Get page number from virtual address
 - Check cache for corresponding frame number
 - If there, use it
 - Checking is much faster than a main memory access
 - If not there, access memory to get frame number
 - And load it into the cache
 - Add page offset to frame address

Hit Ratio

- Percent of time page number is found in cache
- Used to measure efficiency of caching
- Example: 50ns to search cache, 750ns to access memory
 - In cache: access time is $50\text{ns} + 750\text{ns} = 800\text{ns}$
 - Not in cache: access time is $50\text{ns} + 750\text{ns} + 750\text{ns} = 1550\text{ns}$

Effective Memory Access Time

- Average time needed for a memory reference:

hit ratio \times time needed to reference page when page number in cache
+
(1 – hit ratio) \times time needed to reference page when page number not
in cache

Examples

- Building on the earlier one:
- 80% hit ratio: EMAT is $0.8 \times 800\text{ns} + (1-0.8) \times 1550\text{ns} = 956\text{ns}$
 - Slowdown is $(956 - 750) / 750 = 0.274 = 27.4\%$
- 90% hit ratio: EMAT is $0.9 \times 800\text{ns} + (1-0.9) \times 1550\text{ns} = 875\text{ns}$
 - Slowdown is $(875 - 750) / 750 = 0.167 = 16.7\%$

Sharing Pages

- Re-entrant code: code that is not altered
 - Also called pure code, non-self modifying code
- Just put appropriate entries in page table!
- Example: program instructions take up 250 pages, data at most 200 pages
 - With sharing: $250 + 200 + 200 = 650$ frames used
 - Without sharing: $250 + 200 + 250 + 200 = 900$ frames used
- Note: critical that shared pages not be altered!
 - This means the operating system must enforce this

Protection

- Protection bits associated with each page
 - Kept in the page table
- 1 bit to indicate if page is read/write or read only
- 1 bit to indicate whether value in page table is valid or invalid
- More bits for other forms of protection
- So during computation of physical address, operating system can verify the access is appropriate
 - If not (writing a read only page, accessing an invalid entry), trap to operating system

Trapping Illegal Addresses

- Uses the bits to allow or disallow access
- Example:
 - Page size 2048 words per page
 - Program uses addresses 0 . . 10040 (5 pages)
- Suppose it tries to access page 6
 - That's memory address 12288, which is not in program's address space
- Trap to operating system!

Example of Fragmentation

- Page size 2048 words per page
- Program uses addresses 0 . . 10068 (5 pages)
- 5 pages uses 10240 addresses
- So internal fragmentation is 200 words (space left over in page 5)
 - Cannot deny access to those words as you can't block access to specific *words*, just pages
 - It's all of a page or no part of a page

Alternate View of Memory

- User view: program sees memory as a contiguous memory space
 - The memory is divided into equally-sized blocks of instructions or data (pages)
- OS view: OS sees user's program scattered throughout physical memory
- How do we reconcile these?

Reconciliation

- Address translation mechanism maps virtual memory locations to physical locations under control of operating systems
- So physical and virtual addresses may be different
- Example: XNS-940 had virtual address space of 14 bits, but physical address space of 16 bits
 - Page number (3 bits) referenced page table entry to get 5 bit frame number
 - So 4 times as much physical memory as virtual memory
- Widely used when address spaces grow
 - Example: 16 bit address space grows to 32 bit address space

Reconciliation

- Widely used when address spaces grow
 - Example: 16 bit address space grows to 32 bit address space
 - Virtual addresses still 16 bits
 - Physical addresses become 32 bits
- Can't use more memory than before

Tracking Used Frames

- Operating system keeps track of what frames are used and which are not, total number of frames, etc.
- Stored in a global frame table
 - Like a page table, but has one entry per *frame*, not per page
 - Entries indicate if frames are allocated and, if so, to which process

Segmentation

- View program as collection of variable-sized segments
 - 1 segment per function or data structure
 - Segments are of variable length
 - Words identified by offsets into segments
- Called *segmentation*
 - Virtual address space is collection of segments
 - Segments have name and length
 - Addresses specify name of segment, offset into that segment

Segment Names

- These are numbers
 - It's the easiest thing to do
- Segments often generated by compiler — look for something like “.text n ”, which says what follows is in segment number n
- Example: C program may have:
 - segment for global variables
 - segment for program stack
 - segments for instructions of each function
 - segments for local variables of each function

How Segmentation Is Done

- Associated with each segment table are 2 registers
 - Segment table base register (STBR) holds address of start of segment table
 - Segment table limit register (STLR) holds highest address of segment table
- Addresses are (s, d)
 - s is segment number
 - d is offset into segment

How Segmentation Is Done

- Add s to value stored in STBR
- If it exceeds value in STLR, trap; it's an illegal address reference
- Use value to get segment table entry
 - This has (*limit*, *base*) physical addresses
- Compare limit to d
- If d exceeds *limit*, trap; it's an illegal address reference
- Add d to *limit* to get physical memory address

How Segmentation Is Done

